

**SIT330-770: Natural Language Processing**

**Week 3 - Text processing**

Regular Expressions, Text Normalization, Edit Distance

**Dr. Mohamed Reda Bouadjene**

School of Information Technology, Faculty of Sci Eng & Built Env

[reda.bouadjene@deakin.edu.au](mailto:reda.bouadjene@deakin.edu.au)



DEAKIN UNIVERSITY


1

**SIT330-770: Natural Language Processing**

**Week 3.1 - Regular Expressions**

**Dr. Mohamed Reda Bouadjene**

School of Information Technology, Faculty of Sci Eng & Built Env




DEAKIN UNIVERSITY

2

**Regular expressions**

- A formal language for specifying text strings
- How can we search for any of these?
  - woodchuck
  - woodchucks
  - Woodchuck
  - Woodchucks



DEAKIN UNIVERSITY

3

**Regular Expressions: Disjunctions**

- Letters inside square brackets []
 

Pattern	Matches
[wW]oodchuck	Woodchuck, woodchuck
[1234567890]	Any digit
- Ranges [A-Z]
 

Pattern	Matches
[A-Z]	An upper case letter
[a-z]	A lower case letter
[0-9]	A single digit

DEAKIN UNIVERSITY

4

**Regular Expressions: Negation in Disjunction**

- Negations [^ss]
  - Carat means negation only when first in []

Pattern	Matches
[^A-Z]	Not an upper case letter
[^Ss]	Neither 'S' nor 's'
[^e^]	Neither e nor ^
a^b	The pattern a carat b


DEAKIN UNIVERSITY

5

**Regular Expressions: More Disjunction**

- Woodchuck is another name for groundhog!
- The pipe | for disjunction

Pattern	Matches
groundhog woodchuck	woodchuck
yours mine	yours
a b c	= [abc]
[gG]roundhog [wW]oodchuck	Woodchuck



DEAKIN UNIVERSITY

6

Regular Expressions: **? \* + .**

Pattern	Matches
<code>colou?r</code>	Optional previous char <u>color</u> <u>colour</u>
<code>oo*h!</code>	0 or more of previous char <u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
<code>o+h!</code>	1 or more of previous char <u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
<code>baa+</code>	<u>baa</u> <u>baaa</u> <u>baaaa</u> <u>baaaaa</u>
<code>beg.n</code>	<u>begin</u> <u>begun</u> <u>begun</u> <u>beg3n</u>

Stephen C Kleene  
Kleene \*, Kleene +

7

7

Regular Expressions: Anchors **^ \$**

Pattern	Matches
<code>^[A-Z]</code>	<u>P</u> alo Alto
<code>^[^A-Za-z]</code>	<u>"</u> 1 "Hello"
<code>\.\$</code>	The end <u>.</u>
<code>.\$</code>	The end <u>?</u> The end! <u>.</u>

8

8

Example

- Find me all instances of the word "the" in a text.

the

Misses capitalized examples

[tT]he

Incorrectly returns other or theology

[^a-zA-Z][tT]he[^a-zA-Z]

9

9

Errors

- The process we just went through was based on fixing two kinds of errors:
  - Matching strings that we should not have matched (there, then, other)  
**False positives (Type I errors)**
  - Not matching things that we should have matched (The)  
**False negatives (Type II errors)**

10

10

Errors cont.

- In NLP we are always dealing with these kinds of errors.
- Reducing the error rate for an application often involves two antagonistic efforts:
  - Increasing accuracy or precision (minimizing false positives)
  - Increasing coverage or recall (minimizing false negatives).

11

11

Summary

- Regular expressions play a surprisingly large role
  - Sophisticated sequences of regular expressions are often the first model for any text processing
- For hard tasks, we use machine learning classifiers
  - But regular expressions are still used for pre-processing, or as features in the classifiers
  - Can be very useful in capturing generalizations

12

12

## SIT330-770: Natural Language Processing

Week 3.2- More Regular Expressions: Substitutions and ELIZA

Dr. Mohamed Reda Bouadjeneq

School of Information Technology,  
Faculty of Sci Eng & Built Env



13

## Substitutions

- Substitution in Python and UNIX commands:
 

```
s/regexp1/pattern/
```

 e.g.:
 

```
s/colour/color/
```

14

## Capture Groups

- Say we want to put angles around all numbers:
 

```
the 35 boxes → the <35> boxes
```
- Use parens () to "capture" a pattern into a numbered register (1, 2, 3...)
- Use \1 to refer to the contents of the register
 

```
s/([0-9]+)/<\1>/
```

15

## Capture groups: multiple registers

- /the (.\*?)er they (.\*), the \1er we \2/
- Matches
 

```
the faster they ran, the faster we ran
```
- But not
 

```
the faster they ran, the faster we ate
```

16

## But suppose we don't want to capture?

- Parentheses have a double function: grouping terms, and capturing
- Non-capturing groups: add a ?: after paren:
- E.g.: /(?:some|a few) (people|cats) like some \1/
  - matches
    - some cats like some cats
  - but not
    - some cats like some some

17

## Lookahead assertions

- (?= pattern) is true if pattern matches, but is **zero-width**; **doesn't advance character pointer**
- (?! pattern) true if a pattern does not match
- How to match, at the beginning of a line, any single word that doesn't start with "Volcano":
 

```
o/^(?!Volcano)[A-Za-z]+/
```

18

Simple Application: ELIZA

- Early NLP system that imitated a Rogerian psychotherapist
  - Joseph Weizenbaum, 1966.
- Uses pattern matching to match, e.g.,:
  - "I need X"
 and translates them into, e.g.
  - "What would it mean to you if you got X?"

19

19

Simple Application: ELIZA

Men are all alike.  
 IN WHAT WAY  
 They're always bugging us about something or other. CAN YOU THINK OF A SPECIFIC EXAMPLE  
 Well, my boyfriend made me come here.  
 YOUR BOYFRIEND MADE YOU COME HERE  
 He says I'm depressed much of the time.  
 I AM SORRY TO HEAR YOU ARE DEPRESSED

20

20

How ELIZA works

- s/. \* I'M (depressed|sad) . \*/ I AM SORRY TO HEAR YOU ARE \1/
- s/. \* I AM (depressed|sad) . \*/ WHY DO YOU THINK YOU ARE \1/
- s/. \* all . \*/ IN WHAT WAY?/
- s/. \* always . \*/ CAN YOU THINK OF A SPECIFIC EXAMPLE?/

21

21

SIT330-770: Natural Language Processing

Week 3.3 - Words and Corpora

Dr. Mohamed Reda Bouadjene

School of Information Technology,  
 Faculty of Sci Eng & Built Env

22

22

How many words in a sentence?

- "I do uh main- mainly business data processing"
- Fragments, filled pauses
- "Seuss's **cat** in the hat is different from other **cats**!"
- **Lemma**: same stem, part of speech, rough word sense
  - **cat** and **cats** = same lemma
- **Wordform**: the full inflected surface form
  - **cat** and **cats** = different wordforms

23

23

How many words in a sentence?

they lay back on the San Francisco grass and looked at the stars and their

- **Type**: an element of the vocabulary.
- **Token**: an instance of that type in running text.
- How many?
  - 15 tokens (or 14)
  - 13 types (or 12) (or 11?)

24

24

### How many words in a corpus?

**N** = number of tokens  
**V** = vocabulary = set of types, **|V|** is size of vocabulary

Heaps Law = Herdan's Law =  $|V| = kN^\beta$ , where often  $0.67 < \beta < 0.75$   
 i.e., vocabulary size grows with > square root of the number of word tokens

	Tokens = N	Types =  V
Switchboard phone conversations	2.4 million	20 thousand
Shakespeare	884,000	31 thousand
COCA	440 million	2 million
Google N-grams	1 trillion	13+ million

25

### Corpora

Words don't appear out of nowhere!  
 A text is produced by

- a specific writer(s),
- at a specific time,
- in a specific variety,
- of a specific language,
- for a specific function.

26

### Corpora vary along dimension like

- **Language:** 7097 languages in the world
- **Variety**, like African American Language varieties.
  - AAE Twitter posts might include forms like "iont" (*I don't*)
- **Code switching**, e.g., Spanish/English, Hindi/English:
 

S/E: Por primera vez veo a @username actually being hateful! It was beautiful:  
 [For the first time I get to see @username actually being hateful! It was beautiful.]

H/E: dost tha or ra- hega ... don't worry ... but dherya rakhe  
 ["he was and will remain a friend ... don't worry ... but have faith"]
- **Genre:** newswire, fiction, scientific articles, Wikipedia
- **Author Demographics:** writer's age, gender, ethnicity, SES

27

### Corpus datasheets

Gebru et al (2020), Bender and Friedman (2018)

**Motivation:**

- Why was the corpus collected?
- By whom?
- Who funded it?

**Situation:** In what situation was the text written?

**Collection process:** If it is a subsample how was it sampled? Was there consent? Pre-processing?

- +Annotation process, language variety, demographics, etc.

28

### SIT330-770: Natural Language Processing

Week 3.4 - Word tokenization

Dr. Mohamed Reda Bouadjene

School of Information Technology,  
 Faculty of Sci Eng & Built Env

29

### Text Normalization

- Every NLP task requires text normalization:
  1. Tokenizing (segmenting) words
  2. Normalizing word formats
  3. Segmenting sentences

30

### Space-based tokenization

- A very simple way to tokenize
  - For languages that use space characters between words
    - Arabic, Cyrillic, Greek, Latin, etc., based writing systems
  - Segment off a token between instances of spaces
- Unix tools for space-based tokenization
  - The "tr" command
  - Inspired by Ken Church's UNIX for Poets
  - Given a text file, output the word tokens and their frequencies

31

### Simple Tokenization in UNIX

- Given a text file, output the word tokens and their frequencies

```
tr -sc 'A-Za-z' '\n' < shakes.txt | sort | uniq -c
```

Change all non-alpha to newlines  
Sort in alphabetical order  
Merge and count each type

```
1945 A
 72 AARON
 19 ABESS
  5 ABBOT
... ..
```

32

### The first step: tokenizing

```
tr -sc 'A-Za-z' '\n' < shakes.txt | head
```

```
THE
SONNETS
by
William
Shakespeare
From
fairrest
creatures
We
...
```

33

### The second step: sorting

```
tr -sc 'A-Za-z' '\n' < shakes.txt | sort | head
```

```
A
A
A
A
A
A
A
A
...
```

34

### More counting

- Merging upper and lower case
 

```
tr 'A-Z' 'a-z' < shakes.txt | tr -sc 'A-Za-z' '\n' | sort | uniq -c
```
- Sorting the counts
 

```
tr 'A-Z' 'a-z' < shakes.txt | tr -sc 'A-Za-z' '\n' | sort | uniq -c | sort -n -r
```

```
23243 the
22225 i
18618 and
16339 to
15687 of
12780 a
12163 you
10839 my
10005 in
 8954 d
```

What happened here?

35

### Issues in Tokenization

- Can't just blindly remove punctuation:
  - m.p.h., Ph.D., AT&T, cap'n
  - prices (\$45.55)
  - dates (01/02/06)
  - URLs (<http://www.stanford.edu>)
  - hashtags (#nlproc)
  - email addresses (someone@cs.colorado.edu)
- Clitic: a word that doesn't stand on its own
  - "are" in *we're*, French "je" in *j'ai*, "le" in *l'honneur*
- When should multiword expressions (MWE) be words?
  - New York, rock 'n' roll

36

### Tokenization in NLTK

```
>>> text = 'That U.S.A. poster-print costs $12.40...'
>>> pattern = r'((?x) # set flag to allow verbose regexps
... ([A-Z]\.)* # abbreviations, e.g. U.S.A.
... | \w+(-\w+)* # words with optional internal hyphens
... | \$?(d+\.d+)?%? # currency and percentages, e.g. $12.40, 82%
... | \.\.\. # ellipsis
... | [\'\".,;:?"'()@-_-]' # these are separate tokens; includes ], [
... )'
>>> nltk.regexp_tokenize(text, pattern)
['That', 'U.S.A.', 'poster-print', 'costs', '$12.40', '...']
```

37

### Tokenization in NLTK

```
>>> text = 'That U.S.A. poster-print costs $12.40...'
>>> pattern = r'((?x) # set flag to allow verbose regexps
... ([A-Z]\.)* # abbreviations, e.g. U.S.A.
```

38

### Tokenization in languages without spaces

- Many languages (like Chinese, Japanese, Thai) don't use spaces to separate words!
- How do we decide where the token boundaries should be?

39

### Word tokenization in Chinese

- Chinese words are composed of characters called "**hanzi**" (or sometimes just "**zi**")
- Each one represents a meaning unit called a morpheme.
- Each word has on average 2.4 of them.
- But deciding what counts as a word is complex and not agreed upon.

40

### How to do word tokenization in Chinese?

姚明进入总决赛 "Yao Ming reaches the finals"

- 3 words?

姚明 进入 总决赛

YaoMing reaches finals

- 5 words?

姚 明 进 入 总 决 赛

Yao Ming reaches overall finals

- 7 characters? (don't use words at all):

姚 明 进 入 总 决 赛

Yao Ming enter enter overall decision game

41

### Word tokenization / segmentation

- So, in Chinese it's common to just treat each character (zi) as a token.
  - So, the **segmentation** step is very simple
- In other languages (like Thai and Japanese), more complex word segmentation is required.
  - The standard algorithms are neural sequence models trained by supervised machine learning.

42

## SIT330-770: Natural Language Processing

Week 3.5 - Byte Pair Encoding

Dr. Mohamed Reda Bouadjene

School of Information Technology,  
Faculty of Sci Eng & Built Env



43

### Another option for text tokenization

- Instead of
  - white-space segmentation
  - single-character segmentation
- **Use the data** to tell us how to tokenize.
- **Subword tokenization** (because tokens can be parts of words as well as whole words)

44

### Subword tokenization

- Three common algorithms:
  - **Byte-Pair Encoding (BPE)** (Sennrich et al., 2016)
  - **Unigram language modeling tokenization** (Kudo, 2018)
  - **WordPiece** (Schuster and Nakajima, 2012)
- All have 2 parts:
  - A token **learner** that takes a raw training corpus and induces a vocabulary (a set of tokens).
  - A token **segmenter** that takes a raw test sentence and tokenizes it according to that vocabulary

45

### Byte Pair Encoding (BPE) token learner

- Let vocabulary be the set of all individual characters  
= {A, B, C, D, ..., a, b, c, d, ...}
- Repeat:
  - Choose the two symbols that are most frequently adjacent in the training corpus (say 'A', 'B')
  - Add a new merged symbol 'AB' to the vocabulary
  - Replace every adjacent 'A' 'B' in the corpus with 'AB'.
- Until  $k$  merges have been done.

46

### BPE token learner algorithm

```

function BYTE-PAIR ENCODING(strings  $C$ , number of merges  $k$ ) returns vocab  $V$ 
 $V \leftarrow$  all unique characters in  $C$       # initial set of tokens is characters
for  $i = 1$  to  $k$  do                      # merge tokens til  $k$  times
   $t_L, t_R \leftarrow$  Most frequent pair of adjacent tokens in  $C$ 
   $t_{NEW} \leftarrow t_L + t_R$           # make new token by concatenating
   $V \leftarrow V + t_{NEW}$               # update the vocabulary
  Replace each occurrence of  $t_L, t_R$  in  $C$  with  $t_{NEW}$  # and update the corpus
return  $V$ 

```

47

### Byte Pair Encoding (BPE) Addendum

- Most subword algorithms are run inside space-separated tokens.
- So we commonly first add a special end-of-word symbol '\_\_\_' before space in training corpus
- Next, separate into letters.

48



**BPE token learner**

- Original (very fascinating 😊) corpus:
- low low low low low lowest lowest newer newer newer newer newer
- newer wider wider wider new new
- Add end-of-word tokens, resulting in this vocabulary:

Corpus	Vocabulary
5 low _	_, d, e, i, l, n, o, r, s, t, w
2 low est _	
6 new er _	
3 wid er _	
2 new _	

49

49

**BPE**

Corpus	Vocabulary
5 low _	_, d, e, i, l, n, o, r, s, t, w
2 low est _	
6 new er _	
3 wid er _	
2 new _	

**Merge [e r] to [er]**

Corpus	Vocabulary
5 low _	_, d, e, i, l, n, o, r, s, t, w, <b>er</b>
2 low est _	
6 new <b>er</b> _	
3 wid <b>er</b> _	
2 new _	

50

50

**BPE**

Corpus	Vocabulary
5 low _	_, d, e, i, l, n, o, r, s, t, w, er
2 low est _	
6 new er _	
3 wid er _	
2 new _	

**Merge [er \_] to [er\_]**

Corpus	Vocabulary
5 low _	_, d, e, i, l, n, o, r, s, t, w, er, <b>er_</b>
2 low est _	
6 new <b>er_</b> _	
3 wid <b>er_</b> _	
2 new _	

51

51

**BPE**

Corpus	Vocabulary
5 low _	_, d, e, i, l, n, o, r, s, t, w, er, er_
2 low est _	
6 new er _	
3 wid er _	
2 new _	

**Merge [n e] to [ne]**

Corpus	Vocabulary
5 low _	_, d, e, i, l, n, o, r, s, t, w, er, er_, <b>ne</b>
2 low est _	
6 <b>ne</b> w er _	
3 wid er _	
2 <b>ne</b> w _	

52

52

**BPE**

- The next merges are:

Merge	Current Vocabulary
(ne, w)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new
(l, o)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new, lo
(lo, w)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new, lo, low
(new, er_)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new, lo, low, newer_
(low, _)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new, lo, low, newer_, low_

53

53

**BPE token segmenter algorithm**

On the test data, run each merge learned from the training data:

- Greedily
- In the order we learned them
- (test frequencies don't play a role)

So: merge every [e r] to [er], then merge [er \_] to [er\_], etc.

- Result:
  - Test set "n e w e r \_" would be tokenized as a full word
  - Test set "l o w e r \_" would be two tokens: "low er\_"

54

54

Properties of BPE tokens

- Usually include frequent words and frequent subwords
  - Which are often morphemes like *-est* or *-er*
- A **morpheme** is the smallest meaning-bearing unit of a language
  - unlikeliest* has 3 morphemes *un-*, *likely*, and *-est*

55

55

SIT330-770: Natural Language Processing

Week 3.6 - Word Normalization and other issues

Dr. Mohamed Reda Bouadjeneq

School of Information Technology,  
Faculty of Sci Eng & Built Env

56

56

Word Normalization

- Putting words/tokens in a standard format
  - U.S.A. or USA
  - uhhuh or uh-huh
  - Fed or fed
  - am, is, be, are

57

57

Case folding

- Applications like IR: reduce all letters to lower case
  - Since users tend to use lower case
  - Possible exception: upper case in mid-sentence?
    - e.g., *General Motors*
    - Fed* vs. *fed*
    - SAIL* vs. *sail*
- For sentiment analysis, MT, Information extraction
  - Case is helpful (*US* versus *us* is important)

58

58

Lemmatization

- Represent all words as their lemma, their shared root = dictionary headword form:
  - am, are, is* → *be*
  - car, cars, car's, cars'* → *car*
  - Spanish *quiero* ('I want'), *quieres* ('you want')  
→ *querer* 'want'
  - He is reading detective stories*  
→ *He be read detective story*

59

59

Lemmatization is done by Morphological Parsing

- Morphemes:
  - The small meaningful units that make up words
  - Stems**: The core meaning-bearing units
  - Affixes**: Parts that adhere to stems, often with grammatical functions
- Morphological Parsers:
  - Parse *cats* into two morphemes *cat* and *s*
  - Parse Spanish *amaren* ('if in the future they would love') into morpheme *amar* 'to love', and the morphological features *3PL* and *future subjunctive*.

60

60

### Stemming

- Reduce terms to stems, chopping off affixes crudely

This was not the map we found in Billy Bones's chest, but an accurate copy, complete in all things-names and heights and soundings-with the single exception of the red crosses and the written notes.

→

Thi wa not the map we found in Billi Bone s chest but an accur copi complet in all thing name and height and sound with the singl except of the red cross and the written note.

61

### Porter Stemmer

- Based on a series of rewrite rules run in series
  - A cascade, in which output of each pass fed to next pass
- Some sample rules:
 

ATIONAL	→	ATE (e.g., ATIONAL→ATE)
ING	→	ε if stem contains vowel (e.g., motoring → motor)
SSSES	→	SS (e.g., grasses → grass)

62

### Dealing with complex morphology is necessary for many languages

- e.g., the Turkish word:
- Uygarlastiramadiklarimizdanmissinizcasina
- '(behaving) as if you are among those whom we could not civilize'
- Uygar 'civilized' + las 'become'
  - + tir 'cause' + ama 'not able'
  - + dik 'past' + lar 'plural'
  - + imiz 'p1pl' + dan 'abl'
  - + mis 'past' + siniz '2pl' + casina 'as if'

63

### Sentence Segmentation

- ! , ? mostly unambiguous but period "." is very ambiguous
  - Sentence boundary
  - Abbreviations like Inc. or Dr.
  - Numbers like .02% or 4-3
- Common algorithm: Tokenize first: use rules or ML to classify a period as either (a) part of the word or (b) a sentence-boundary.
  - An abbreviation dictionary can help
- Sentence segmentation can then often be done by rules based on this tokenization.

64

### SIT770: Natural Language Processing

Week 3.7 - Definition of Minimum Edit Distance

Dr. Mohamed Reda Bouadjene

School of Information Technology,  
Faculty of Sci Eng & Built Env

65

### How similar are two strings?

- Spell correction
  - The user typed "graffe"
  - Which is closest?
    - graf
    - graft
    - grail
    - giraffe
- Computational Biology
  - Align two sequences of nucleotides
 

```
AGGCTATCACTGACCTGACGCGGATGCCC
TAGCTATCAACGACGCGGATGATTTGCCCGAC
```
  - Resulting alignment:
 

```
--AGGCTATCACTGACCTGACGCGGATGCCC--
TAGCTATCACTGACCTGACGCGGATGATTTGCCCGAC
```

Also for Machine Translation, Information Extraction, Speech Recognition

66

### Edit Distance

- The minimum edit distance between two strings
- Is the minimum number of editing operations
  - Insertion
  - Deletion
  - Substitution
- Needed to transform one into the other

67

### Minimum Edit Distance

- Two strings and their **alignment**:

```

      I N T E * N T I O N
      | | | | | | | |
    * E X E C U T I O N
  
```

68

### Minimum Edit Distance

```

      I N T E * N T I O N
      | | | | | | | |
    * E X E C U T I O N
      d s s i s
  
```

- If each operation has cost of 1
  - Distance between these is 5
- If substitutions cost 2 (Levenshtein)
  - Distance between them is 8

69

### Alignment in Computational Biology

- Given a sequence of bases
 

```

AGGCTATCACCTGACCTCCAGGCCGATGCC
TAGCTATCACGACCGCGGTCGATTGCCCGAC
      
```
- An alignment:
 

```

-AGGCTATCACCTGACCTCCAGGCCGA--TGCCC--
TAG-CTATCAC--GACCGC--GGTCGATTGCCCGAC
      
```
- Given two sequences, align each letter to a letter or gap

70

### Other uses of Edit Distance in NLP

- Evaluating Machine Translation and speech recognition
 

```

R Spokesman confirms senior government adviser was appointed
H Spokesman said the senior adviser was appointed
S I D I
      
```
- Named Entity Extraction and Entity Coreference
  - IBM Inc. announced today
  - IBM profits
  - Stanford Professor Jennifer Eberhardt announced yesterday
  - for Professor Eberhardt...

71

### How to find the Min Edit Distance?

- Searching for a path (sequence of edits) from the start string to the final string:
  - Initial state**: the word we're transforming
  - Operators**: insert, delete, substitute
  - Goal state**: the word we're trying to get to
  - Path cost**: what we want to minimize: the number of edits

```

      intention
     /  |  \
  Del  Ins  Sub
   /    |    \
intention eintention entention
  
```

72

### Minimum Edit as Search

- But the space of all edit sequences is huge!
  - We can't afford to navigate naively
  - Lots of distinct paths wind up at the same state.
    - We don't have to keep track of all of them
    - Just the shortest path to each of those revisited states.

73

73

### Defining Min Edit Distance

- For two strings
  - X of length  $n$
  - Y of length  $m$
- We define  $D(i,j)$ 
  - the edit distance between  $X[1..i]$  and  $Y[1..j]$ 
    - i.e., the first  $i$  characters of X and the first  $j$  characters of Y
  - The edit distance between X and Y is thus  $D(n,m)$

74

74

### SIT770: Natural Language Processing

Week 3.8 - Computing Minimum Edit Distance

Dr. Mohamed Reda Bouadjene

School of Information Technology,  
Faculty of Sci Eng & Built Env

75

75

### Dynamic Programming for Minimum Edit Distance

- Dynamic programming:** A tabular computation of  $D(n,m)$
- Solving problems by combining solutions to subproblems.
- Bottom-up**
  - We compute  $D(i,j)$  for small  $i,j$
  - And compute larger  $D(i,j)$  based on previously computed smaller values
  - i.e., compute  $D(i,j)$  for all  $i$  ( $0 < i < n$ ) and  $j$  ( $0 < j < m$ )

76

76

### Defining Min Edit Distance (Levenshtein)

Initialization

$$D(1,0) = 1$$

$$D(0,j) = j$$

Recurrence Relation:

For each  $i = 1..M$   
For each  $j = 1..N$

$$D(i,j) = \min \begin{cases} D(i-1,j) + 1 \\ D(i,j-1) + 1 \\ D(i-1,j-1) + \begin{cases} 2, & \text{if } X(i) \neq Y(j) \\ 0, & \text{if } X(i) = Y(j) \end{cases} \end{cases}$$

Termination:

$D(N,M)$  is distance

77

77

### The Edit Distance Table

N	9								
O	8								
I	7								
T	6								
N	5								
E	4								
T	3								
N	2								
I	1								
#		1	2	3	4	5	6	7	8
#	E	X	E	C	U	T	I	O	N

78

78

The Edit Distance Table

$$D(i,j) = \min \begin{cases} D(i-1,j) + 1 \\ D(i,j-1) + 1 \\ D(i-1,j-1) + \begin{cases} 2; & \text{if } S_1(i) \neq S_2(j) \\ 0; & \text{if } S_1(i) = S_2(j) \end{cases} \end{cases}$$

N	9									
O	8									
I	7									
T	6									
N	5									
E	4									
T	3									
N	2									
I	1									
#	0	1	2	3	4	5	6	7	8	9
#	E	X	E	C	U	T	I	O	N	

79

79

The Edit Distance Table

$$D(i,j) = \min \begin{cases} D(i-1,j) + 1 \\ D(i,j-1) + 1 \\ D(i-1,j-1) + \begin{cases} 2; & \text{if } S_1(i) \neq S_2(j) \\ 0; & \text{if } S_1(i) = S_2(j) \end{cases} \end{cases}$$

N	9									
O	8									
I	7									
T	6									
N	5									
E	4									
T	3									
N	2									
I	1									
#	0	1	2	3	4	5	6	7	8	9
#	E	X	E	C	U	T	I	O	N	

80

80

The Edit Distance Table

$$D(i,j) = \min \begin{cases} D(i-1,j) + 1 \\ D(i,j-1) + 1 \\ D(i-1,j-1) + \begin{cases} 2; & \text{if } S_1(i) \neq S_2(j) \\ 0; & \text{if } S_1(i) = S_2(j) \end{cases} \end{cases}$$

N	9	8	9	10	11	12	11	10	9	8
O	8	7	8	9	10	11	10	9	8	7
I	7	6	7	8	9	10	9	8	7	6
T	6	5	6	7	8	9	8	7	6	5
N	5	4	5	6	7	8	7	6	5	4
E	4	3	4	5	6	7	6	5	4	3
T	3	2	3	4	5	6	5	4	3	2
N	2	1	2	3	4	5	4	3	2	1
I	1	0	1	2	3	4	3	2	1	0
#	0	1	2	3	4	5	6	7	8	9
#	E	X	E	C	U	T	I	O	N	

81

81

**SIT770: Natural Language Processing**

Week 3.9 - Backtrace for Computing Alignments

Dr. Mohamed Reda Bouadjeneq

School of Information Technology,  
Faculty of Sci Eng & Built Env

82

82

Computing alignments

- Edit distance isn't sufficient
  - We often need to **align** each character of the two strings to each other
- We do this by keeping a "backtrace"
- Every time we enter a cell, remember where we came from
- When we reach the end,
  - Trace back the path from the upper right corner to read off the alignment

83

83

Edit Distance

$$D(i,j) = \min \begin{cases} D(i-1,j) + 1 \\ D(i,j-1) + 1 \\ D(i-1,j-1) + \begin{cases} 2; & \text{if } S_1(i) \neq S_2(j) \\ 0; & \text{if } S_1(i) = S_2(j) \end{cases} \end{cases}$$

N	9									
O	8									
I	7									
T	6									
N	5									
E	4									
T	3									
N	2									
I	1									
#	0	1	2	3	4	5	6	7	8	9
#	E	X	E	C	U	T	I	O	N	

84

84

MinEdit with Backtrace

n	9	8	7	6	5	4	3	2	1	0	9	8
o	8	7	6	5	4	3	2	1	0	9	8	7
i	7	6	5	4	3	2	1	0	9	8	7	6
t	6	5	4	3	2	1	0	9	8	7	6	5
n	5	4	3	2	1	0	9	8	7	6	5	4
e	4	3	2	1	0	9	8	7	6	5	4	3
t	3	2	1	0	9	8	7	6	5	4	3	2
n	2	1	0	9	8	7	6	5	4	3	2	1
i	1	0	9	8	7	6	5	4	3	2	1	0
#	0	1	2	3	4	5	6	7	8	9	8	7
#	e	x	e	c	u	t	i	o	n			

85

Adding Backtrace to Minimum Edit Distance

Base conditions:  $D(i, 0) = i$     Termination:  $D(0, j) = j$      $D(N, M)$  is distance


Recurrence Relation:

For each  $i = 1..M$   
For each  $j = 1..N$

$$D(i, j) = \min \begin{cases} D(i-1, j) + 1 & \text{deletion} \\ D(i, j-1) + 1 & \text{insertion} \\ D(i-1, j-1) + \begin{cases} 0 & \text{if } X(i) = Y(j) \\ 1 & \text{if } X(i) \neq Y(j) \end{cases} & \text{substitution} \end{cases}$$

$ptr(i, j) = \begin{cases} \text{LEFT} & \text{insertion} \\ \text{DOWN} & \text{deletion} \\ \text{DIAG} & \text{substitution} \end{cases}$

86



87

Result of Backtrace

- Two strings and their alignment:

```

I N T E * N T I O N
| | | | |
* E X E C U T I O N

```

88

Performance

- Time:
  - $O(nm)$
- Space:
  - $O(nm)$
- Backtrace
  - $O(n+m)$

89

**SIT770: Natural Language Processing**

Week 3.10 - Weighted Minimum Edit Distance

Dr. Mohamed Reda Bouadjene

School of Information Technology,  
Faculty of Sci Eng & Built Env

90

## Weighted Edit Distance

- Why would we add weights to the computation?

- Spell Correction: some letters are more likely to be mistyped than others
- Biology: certain kinds of deletions or insertions are more likely than others

91

## Confusion matrix for spelling errors

		sub(X, Y) = Substitution of X (incorrect) for Y (correct)																									
		X \ Y																									
X \ Y	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z	
a	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
b	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
c	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
d	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
e	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
f	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
g	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
h	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
i	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
j	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
k	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
l	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
m	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
n	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
o	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
p	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
q	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
r	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
s	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
t	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
u	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
v	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
w	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
x	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
y	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
z	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	

92



93

## Weighted Min Edit Distance

- Initialization:

$$\begin{aligned}
 D(0,0) &= 0 \\
 D(i,0) &= D(i-1,0) + \text{del}[x(i)]; \quad 1 \leq i \leq N \\
 D(0,j) &= D(0,j-1) + \text{ins}[y(j)]; \quad 1 \leq j \leq M
 \end{aligned}$$

- Recurrence Relation:

$$D(i,j) = \min \begin{cases} D(i-1,j) + \text{del}[x(i)] \\ D(i,j-1) + \text{ins}[y(j)] \\ D(i-1,j-1) + \text{sub}[x(i),y(j)] \end{cases}$$

- Termination:

$D(N,M)$  is distance

94

## Where did the name, dynamic programming, come from?

...The 1950s were not good years for mathematical research. [the] Secretary of Defense ...had a pathological fear and hatred of the word, research... I decided therefore to use the word, "programming". I wanted to get across the idea that this was dynamic, this was multistage... I thought, let's ... take a word that has an absolutely precise meaning, namely **dynamic**... it's impossible to use the word, **dynamic**, in a pejorative sense. Try thinking of some combination that will possibly give it a pejorative meaning. It's impossible. Thus, I thought dynamic programming was a good name. It was something not even a Congressman could object to."

Richard Bellman, "Eye of the Hurricane: an autobiography" 1984.

95

## SIT770: Natural Language Processing

Week 3.11 - Minimum Edit Distance in Computational Biology

Dr. Mohamed Reda Bouadjene

School of Information Technology,  
Faculty of Sci Eng & Built Env

96



### Sequence Alignment

AGGCTATCACCTGACCTCCAGGCCGATGCCC  
TAGCTATCACGACCGCGGTCGATTGCCCCGAC

-AGGCTATCACCTGACCTCCAGGCCGA--TGCCC--  
TAG-CTATCAC--GACCGC--GGTCGATTGCCCCGAC

97

97

### Why sequence alignment?

- Comparing genes or regions from different species
  - to find important regions
  - determine function
  - uncover evolutionary forces
- Assembling fragments to sequence DNA
- Compare individuals to looking for mutations

98

98

### Alignments in two fields

- In Natural Language Processing
  - We generally talk about **distance** (minimized)
  - And **weights**
- In Computational Biology
  - We generally talk about **similarity** (maximized)
  - And **scores**

99

99

### The Needleman-Wunsch Algorithm

- Initialization:
  - $D(i, 0) = -i * d$
  - $D(0, j) = -j * d$
- Recurrence Relation:
 
$$D(i, j) = \min \begin{cases} D(i-1, j) & -d \\ D(i, j-1) & -d \\ D(i-1, j-1) + s[x(i), y(j)] \end{cases}$$
- Termination:
  - $D(N, M)$  is distance

100

100

### The Needleman-Wunsch Matrix

(Note that the origin is at the upper left.)

101

101

### A variant of the basic algorithm:

- Maybe it is OK to have an unlimited # of gaps in the beginning and end:

-----CTATCACCTGACCTCCAGGCCGATGCCCTTCCGGC  
GCGAGTTCATCTATCAC--GACCGC--GGTCG-----

- If so, we don't want to penalize gaps at the ends

102

102

### Different types of overlaps

**Example:**  
2 overlapping "reads" from a sequencing project

**Example:**  
Search for a mouse gene within a human chromosome

103

### The Overlap Detection variant

**Changes:**

1. Initialization  
For all  $i, j$ ,  
 $F(i, 0) = 0$   
 $F(0, j) = 0$
2. Termination  
 $F_{opt} = \max \begin{cases} \max_i F(i, N) \\ \max_j F(M, j) \end{cases}$

104

### The Local Alignment Problem

Given two strings  $X = X_1 \dots X_{M_t}$   
 $Y = Y_1 \dots Y_N$

Find substrings  $x', y'$  whose similarity (optimal global alignment value) is maximum

$x = \text{aaaaccccggggtta}$   
 $y = \text{ttccggggaaccaacc}$

105

### The Smith-Waterman algorithm

**Idea:** Ignore badly aligning regions

Modifications to Needleman-Wunsch:

**Initialization:**  $F(0, j) = 0$   
 $F(i, 0) = 0$

**Iteration:**  $F(i, j) = \max \begin{cases} 0 \\ F(i-1, j) - d \\ F(i, j-1) - d \\ F(i-1, j-1) + s(x_i, y_j) \end{cases}$

106

### The Smith-Waterman algorithm

**Termination:**

1. If we want the **best** local alignment...  
 $F_{opt} = \max_{i,j} F(i, j)$   
Find  $F_{opt}$  and trace back
2. If we want **all** local alignments **scoring**  $> t$   
?? For all  $i, j$  find  $F(i, j) > t$ , and trace back?  
Complicated by overlapping local alignments

107

### Local alignment example

$X = \text{ATCAT}$   
 $Y = \text{ATTATC}$

Let:  
 $m = 1$  (1 point for match)  
 $d = 1$  (-1 point for del/ins/sub)

		A	T	T	A	T	C
	0	0	0	0	0	0	0
A	0						
T	0						
C	0						
A	0						
T	0						

108

Local alignment example

X = ATCAT  
Y = ATTATC

		A	T	A	T	C
	0	0	0	0	0	0
A	0	1	0	0	1	0
T	0	0	2	1	0	2
C	0	0	1	1	0	1
A	0	1	0	0	2	1
T	0	0	2	0	1	3

109

Local alignment example

X = **ATCAT**  
Y = **ATTATC**

		A	T	A	T	C
	0	0	0	0	0	0
A	0	1	0	0	1	0
T	0	0	2	1	0	2
C	0	0	1	1	0	1
A	0	1	0	0	2	1
T	0	0	2	0	1	3

110

Local alignment example

X = **ATCAT**  
Y = ATT**ATC**

		A	T	A	T	C
	0	0	0	0	0	0
A	0	1	0	0	1	0
T	0	0	2	1	0	2
C	0	0	1	1	0	1
A	0	1	0	0	2	1
T	0	0	2	0	1	3

111

109

110

111